SpringOne

Getting Started with Spring Cloud Data Flow

Sabby Anandan (@sabbyanandan)
Product Manager

Covering today ..

Spring Cloud Data Flow

- Overview
- Features / Use-cases

Hands-on Labs

- SCDF on Kubernetes
- Event Streaming
- Cloud-native Batch

Housekeeping ..

- Mix of presentation + walkthrough (30 mins)

- 3 hands-on labs (~15 mins each)

- Questions? raise hand; ask questions in chat

Enterprise Integration

“communication between people, machines, computers, and devices and their efficient coordination”

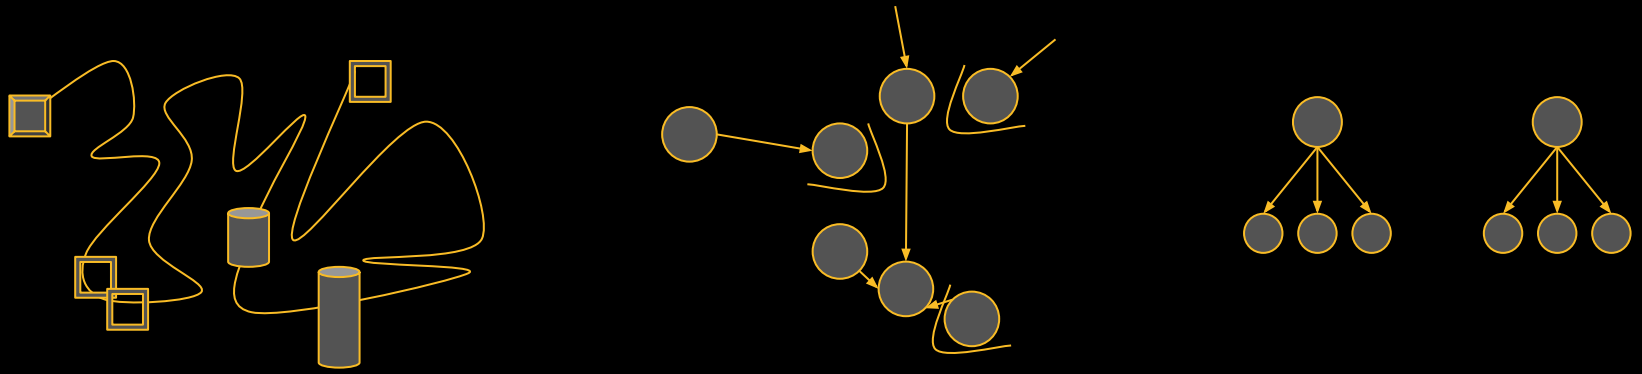

SpringOne

Enterprise Integration Patterns

"communication between people, machines, computers, and devices and their efficient coordination"

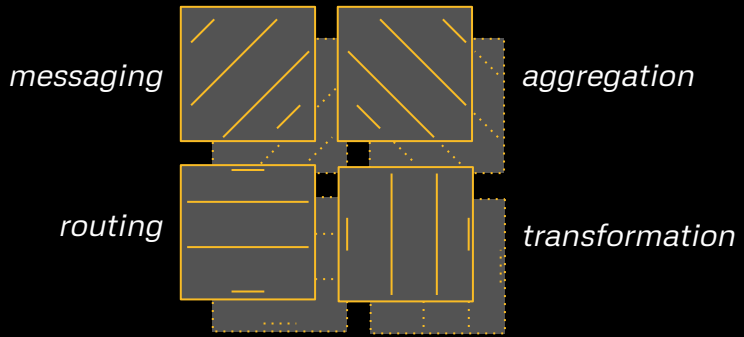

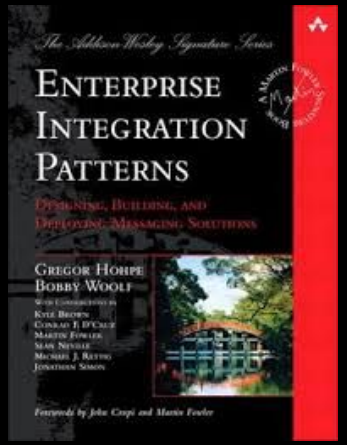

SpringOne

Enterprise Integration in Reality

"communication between people, machines, computers, and devices and their efficient coordination"

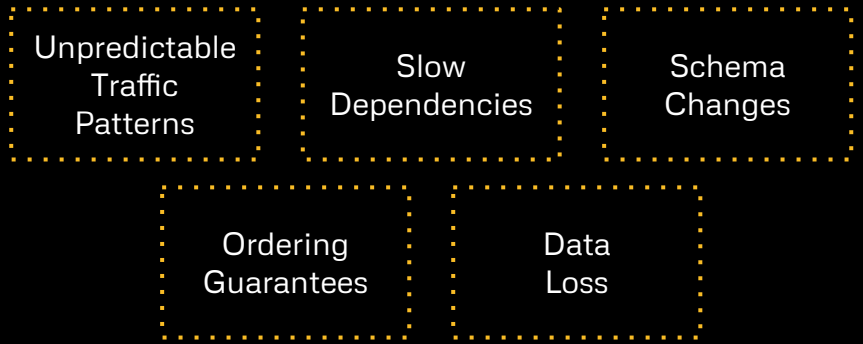

SpringOne

Spring to Rescue!

Cloud-native Enterprise Integration

Spring Cloud Task

Short-lived Spring Boot microservices for batch data processing

Use-cases:
Scheduled data migration jobs
Extract, Transform, and Load (ETL)
Offline machine learning and model training

Spring Cloud Stream

Event-driven Spring Boot microservices for real-time data processing

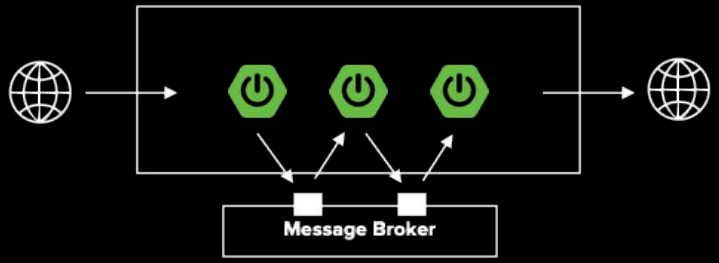

Use-cases:
Enterprise data integration (EAI/EIP)
Event-driven architectures
Real-time predictive analytics (e.g., IoT)

Spring Cloud Data Flow

A microservices based
Streaming and ***Batch***
data processing
in the cloud

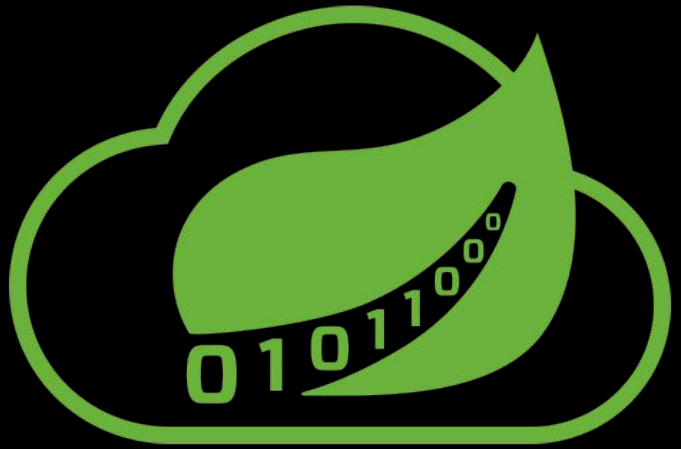

Cloud-native Data Pipelines

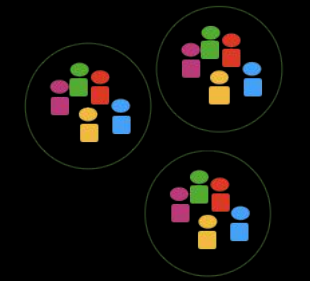

Build Apps

Spring Cloud Stream

Spring Cloud Task

Ship Apps

Deploy

Streams

Monitor

Performance

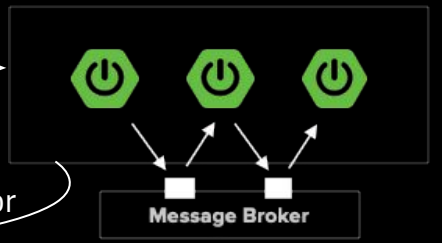

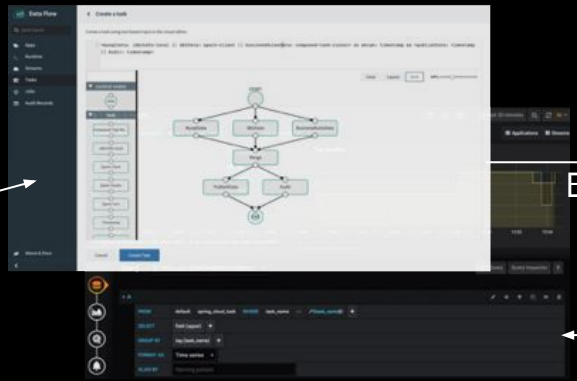

Launch

Batch Jobs

Track

Execution

SpringOne

Spring Cloud Data Flow

Kubernetes

Lab #1: Set up SCDF on Kubernetes

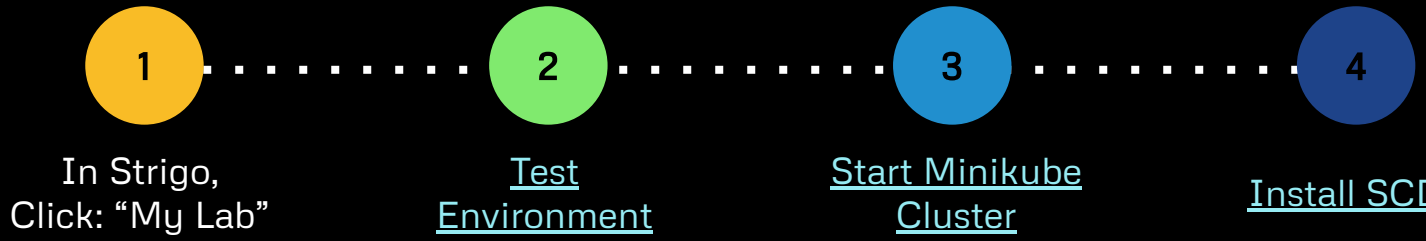
1
2
3
4
In Strigo,
Click: “My Lab”
Test Environment
Start Minikube Cluster
Install SCDF

Cloud-native Data Pipelines

- Extend cloud-native patterns to data processing use-cases
- CI/CD-centered microservices architecture
- Elastically scale Event Streaming and Batch, on-demand

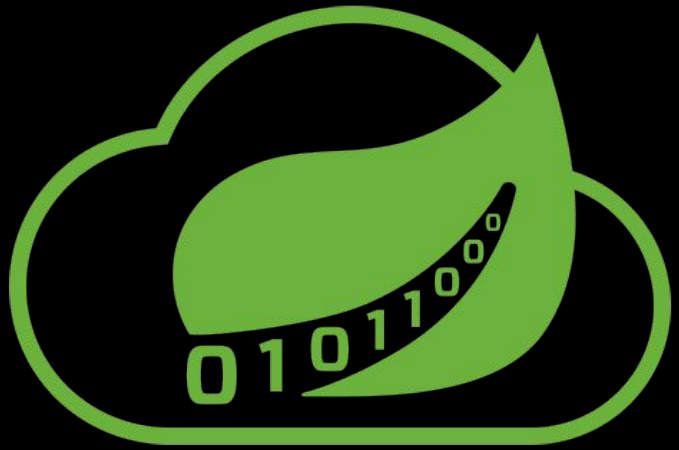

Spring Cloud Data Flow

See: Features

Lab #2: Event Streaming

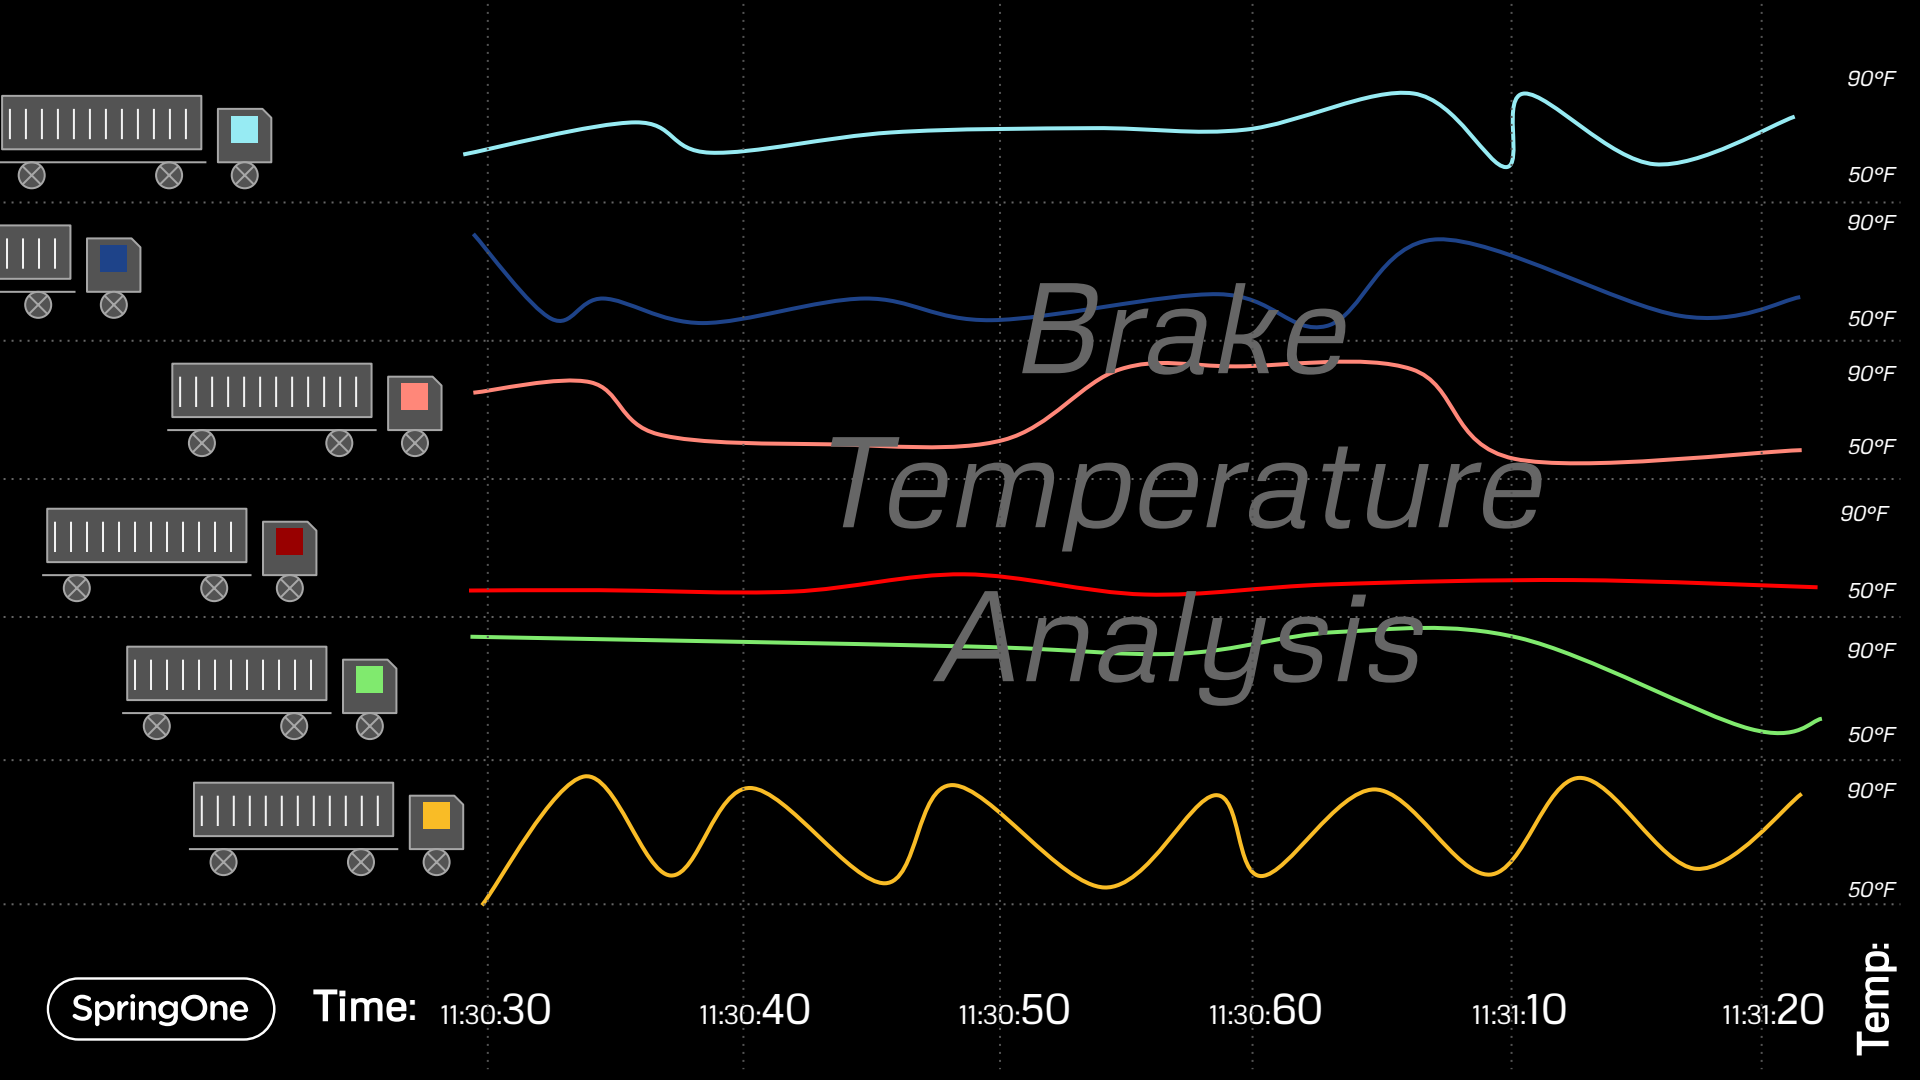
Brake
Temperature
Analysis
90°F
50°F
90°F
50°F
90°F
50°F
90°F
50°F
90°F
50°F
90°F
50°F
SpringOne
Time:
11:30:30
11:30:40
11:30:50
11:30:60
11:31:10
11:31:20
Temp:

```
Supplier<Truck> generateTruck() {
  return () -> randomTruck();
}
```

Topic: trucks

```
Function<?,?> processBrakeTemperature() {
  return input -> input
    .map(..)
    .groupByKey(..)
    .windowedBy(10) // 10-secs moving avg.
    .aggregate(..)
    .map(..);
}
```

Topic: trucks

Moving Window

Avg. Brake Temperature

Topic: truck-logs

```
Consumer<Object> log() {
  return result ->
      System.out.println(result);
}
```

Topic: truck-logs

Results

SpringOne

Lab #2: Cloud-native ETL

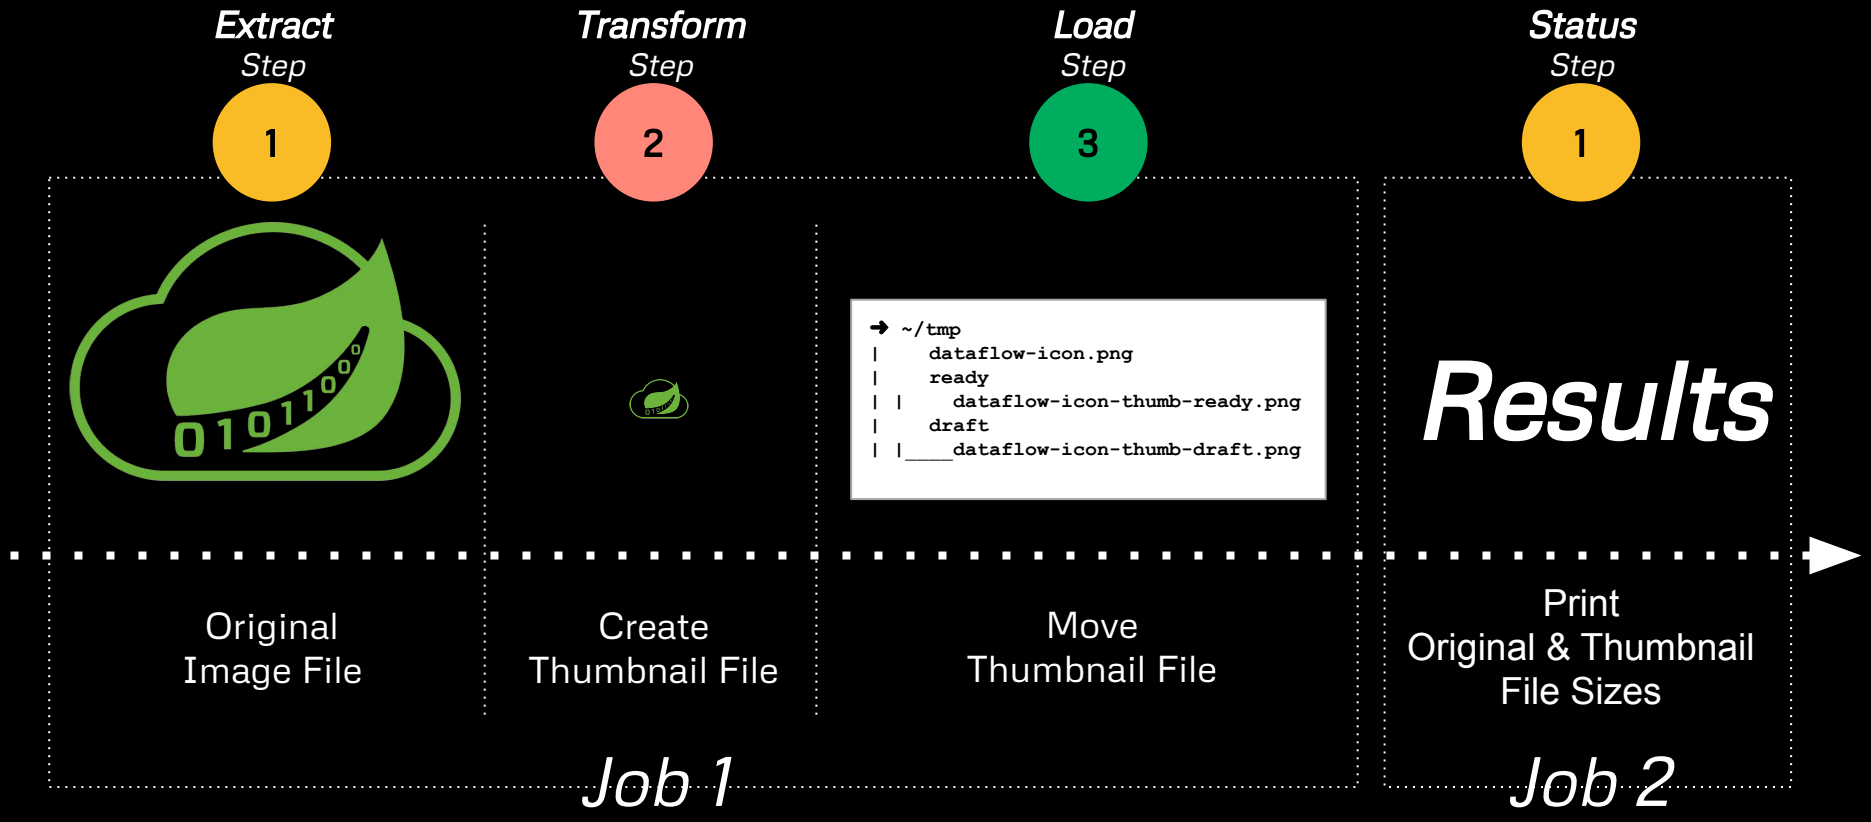
Extract
Step
1
Transform
Step
2
Load
Step
3
Status
Step
1
~/tmp
dataflow-icon.png
ready
dataflow-icon-thumb-ready.png
draft
dataflow-icon-thumb-draft.png
Results
Original
Image File
Create
Thumbnail File
Move
Thumbnail File
Print
Original & Thumbnail
File Sizes
Job 1
Job 2

Resources

Spring Cloud Data Flow: https://dataflow.spring.io/

Source Code: https://github.com/sabbyanandan/SpringOne2020

Lab Instructions: https://hackmd.io/@sabbyanandan/B1bDf74fv

Slides: Speaker Deck

Batching for the Modern Enterprise

Michael Minella, VMware

Sep 2nd

Game of Streams 🐉: How to Tame and Get the Most from Your Messaging Platforms

Mark Heckler, VMware

Sep 2nd

Walking Through Spring Cloud Data Flow

Glenn Renfro, VMware
Ilayaperumal Gopinathan, VMware

Sep 3rd

IoT Scale Event-Stream Processing for Connected Fleet at Penske

Krishna Gogineni, VMware
Shruti Modi, Penske

Sep 3rd